


Pure MLP ladder · training loss

LP pretraining · seed 0 · 2,500 updates per rung · shared axis limits



Each point is ONE batch logged every 250 updates, not a 250-update mean. Source rotation contributes to the zigzags.