

Does joint training preserve source performance?

Fixed validation pairs; x-axis matches source exposure, not total optimizer steps

