










Pure MLP ladder · averaged training curves

Offline W&B rerun · seed 0 · 2,500 updates per rung · shared axis limits



Each point averages all examples in 25 updates (100 points per rung). No additional smoothing.

Source mixtures differ between panels; absolute losses are not a matched-distribution performance comparison.