

Support quality is readout-dependent

Candidate explanation: transfer depends on compatibility between support embeddings and learned class-reference construction, not support quality alone. At 50k, removing context damages simple prototypes but improves native ranking. Fixed-routing value interventions localize the effect; they do not yet predict its sign.

A Support-context utility varies across sources

Suppression minus intact, pooled AUC points; two-stream mean


	Pol.	Elect.	Pages	Bots	Susp.
COVID	-0.1*	+0.0*	+0.1	-4.0	+1.2
Political	-5.7	-0.4	-3.9	-2.4	-2.8
Hong Kong	+9.9	+2.0	+2.5	-7.5	+0.6*
Election	-0.2	-0.0*	-1.3	-10.1	+0.2*
Pages	+1.8	+1.3	+0.9	-1.6	+1.4
Midterm	+3.1	+0.3	+0.8	-0.6	-0.1*
Bots	+0.8	+0.1	+0.1*	-2.6	+0.2*
Ukraine	-0.1	-0.0	-0.0*	-5.4	+1.7
Suspended	+7.4	+0.1*	-0.6	-8.3	+1.4

Teal: suppression helps. Red: it hurts. *: stream signs differ.

Hong Kong is an extreme, not the only affected source. This map establishes heterogeneity, not a validated predictor of its sign.

B Values carry more of the effect than routing

Hong Kong to political; within-episode AUC change (points)



Value-only substitution preserves attention and query vectors exactly. The value-over-key ordering transfers to both 50k streams; baseline/removal endpoints were already known. This is one source-target pair, not cross-source prediction.

C The strongest counterexample: a useful public model needs both roles

Original-style Wiki to FB15K-237, fixed 8,001-update checkpoint; 500 paired 20-way episodes



Both scientific intervals exclude zero:
Support suppression: -33.87 points
95% CI [-34.55, -33.19]
Query suppression: -34.80 points
95% CI [-35.48, -34.14]
The role-conflict prediction fails.

Native exceeds raw by 1.38 points [0.74, 2.03] and untrained by 3.99 [3.36, 4.62]. Intervals are paired episode-bootstrap 95% CIs conditional on this graph/checkpoint, not independent training seeds. Public support/query drift occurs before the metagraph, consistent with shared training-mode BatchNorm; its causal contribution is untested. This is not fixed-query mediation.

D The same supports improve one readout, harm another

Suppression effect: within-episode AUC points, HK to political

Configuration	U1 prototype	Native
2.5k original	+3.31	+7.25
50k original	-6.45	+8.76
50k fresh	-6.35	+10.35

128 episodes/row; fixed U1 queries for prototypes; no new forwards.

Contribution and remaining gap

A localized support-value failure, now distinguished from uniformly better support embeddings. **Not a repair:** at 50k, AUC improves but accuracy falls 49.19% to 25.15%. Next: identify a learned property explaining an affected versus unaffected case. No universal sign rule or deployable method is established.

Verification disclosure: GPU repeats differ numerically. The original logit-tolerance check failed; a dated amendment before baseline/role outcomes required exact labels and all 40,000 predicted classes, which matched. No tolerance was enlarged. Native architecture, sum aggregation, released split/BN behavior and scientific criteria were retained.