

Graph context constructs a classifier

Central distinction: suppressing support context can improve discrimination, or improve ranking while leaving nearly indistinguishable class references and unstable decisions. Similar AUC gains do not establish the same transfer mechanism.

A A real but non-universal phenomenon


Support suppression: pooled AUC change, points

	Pol.	Elect.	Pages	Bots	Susp.
COVID	-0.1*	+0.0*	+0.1	-4.0	+1.2
Political	-5.7	-0.4	-3.9	-2.4	-2.8
Hong Kong	+9.9	+2.0	+2.5	-7.5	+0.6*
Election	-0.2	-0.0*	-1.3	-10.1	+0.2*
Pages	+1.8	+1.3	+0.9	-1.6	+1.4
Midterm	+3.1	+0.3	+0.8	-0.6	-0.1*
Bots	+0.8	+0.1	+0.1*	-2.6	+0.2*
Ukraine	-0.1	-0.0	-0.0*	-5.4	+1.7
Suspended	+7.4	+0.1*	-0.6	-8.3	+1.4

Nine sources, five targets; mean of two episode streams, not training seeds. Teal helps, red hurts; * signs differ. Hong Kong is an extreme, not the only affected source.

B Similar gains, different stability

Suppression gain in within-episode AUC, points




Current initialization Training label table

Early accuracy improves by 7.7-7.9 points under either initialization. Late accuracy worsens under both. Lower early initialization sensitivity was predicted from its less-collapsed references before the new factorial outcomes. **One source-target pair; protocols differ.**

C An explained ablation side effect


Late original: support unit dispersion (log scale)



A large learned MLP(0) first overwhelms feature variation. Subtracting it restores pre-BN diversity, but not downstream geometry; it also harms intact AUC by 3.2/4.2 points. **Offset removal is not a transfer repair.** Fresh-stream behavior agrees.

D Generality result: both roles can help

Public Wiki to FB15K-237: accuracy (%)



A useful original-style checkpoint needs both roles: support/query suppression costs 33.87/34.80 points. 500 paired 20-way episodes, 8,001 updates. This contradicts universal support harm. Public queries are not fixed under the role intervention; do not claim the social mediation result generalizes.

Contribution supported: role-specific interventions plus class-reference geometry distinguish robust discrimination gains from initialization-sensitive ranking gains. The early benefit survives the strongest tested interface explanation; the late ablation has a concrete but incomplete operating-point explanation. **Not yet supported:** a rule predicting when natural support context helps, a broadly useful adaptation method, or an 8+ contribution claim.

Decision: keep the mechanism-paper direction; close offset and initialization tuning. The unresolved scientific target is the early robust benefit, not another repair of the late ablation. No manuscript expansion or new sweep until that distinction guides the next hypothesis.

Evidence: nine-source map; initialization factorials (9e19e5eb); paired offset test (48907a3b); public quality/roles. Public logit-tolerance replay failed; an explicit amendment before scientific outcomes required exact labels and all 40,000 predicted classes, which matched. No tolerances were enlarged. These are conditional results, not seed-level generality.