


What does joint training add beyond averaging specialists?

All foreign-source mixtures; fixed equal-probability ensembles; identical test inputs within each stream.



Top: both gains are relative to mean constituent AUC; above the diagonal favors the trained mixture over the ensemble.

Bottom: below the diagonal means more shared errors are fixed than unanimous-correct decisions are harmed.

The bottom panel is an accuracy decomposition, not an AUC decomposition; mixed-correctness queries are reported in the tables.

Historical sampler; one training seed. Ensembles use 2x/8x training updates and model forwards, not measured wall-clock cost.

Cases share sources and are not independent trials. No fitted ensemble weights, target-selected checkpoints, or causal interference claim. Axis ranges vary by target.