


Specialist averaging across saved training-update budgets

Joint models stay fixed at 2,500 updates. Above zero favors the specialist ensemble; shaded region uses no more total updates.

— Original episodes — Fresh episodes ◆ Fixed under-budget checkpoint rule



Pairs: every one of 28 foreign-source cases per target/stream is shown faintly; bold line is their mean. LOO: one foreign case per target/stream.

Diamonds select the largest saved step with $K \times \text{step} \leq 2,500$: pairs at 900 each (1,800 total), LOO at 300 each (2,400 total).

All four saved checkpoints and both episode streams are retained; no target-based checkpoint or ensemble-weight fitting.

One historical training seed; equal-probability averaging. Saved-update/episode budgets are not matched FLOPs or wall time.

Ensembles retain 2x/8x model capacity and inference forwards. Shared sources are not independent replications; no causal interference claim. Axis ranges vary.