

What changes during training?

AUC at 2,500 updates minus AUC at 100 updates



One historical training seed; fixed inputs within each stream. Includes source=target cells.
Common-decoder changes are not a causal decomposition or proof of information loss.