


SAMGPT native objective: the constituent-specialist max rule is approximate

loss is lower-better, so the max-performance rule is the minimum specialist loss · ladder evaluated on CARC V100; specialist matrix on Tucker H100 · same fixed unseen views

243 target-level comparisons



Observed vs predicted mean curves

