


Native-pretraining mixture ladders → downstream label classification



Each colored line is one common labeled target; black is their unweighted mean. Every model uses its native pretext. VISION has only odd rungs; family-specific fixed-compute endpoints differ.