


SAMGPT fits its native objective nearly to zero at every ladder rung

Loss rises modestly with mixture size, but source composition remains a large factor



200 training epochs per rung · mean of per-source BCE losses · best epoch chosen per run · no validation split