


Native-objective pair impact · targets shown by color

Targets grouped in legend order · separate metric scales · black bars = medians



GraphMAE reconstruction: no per-target native evaluation grid found. Its static-link results are a downstream task and are excluded.

PRODIGY: 2,500 updates / seed 0. GraphSAGE: SSL-validation-selected checkpoints / seed 0. SAMGPT: 500 updates / seed 39.

One training seed per model; includes source-seen and held-out targets. Native tasks, evaluation protocols, and scales differ; dots are dependent.