


Recent pair experiments · impact by model and objective

36 source pairs × 5 targets per model · target colors · black bars = medians



360 dots/model on the left; 180 on the right. One training seed per model; contrasts share models and are not independent.

PRODIGY: 2,500 updates, seed 0. GraphSAGE: SSL-validation-selected checkpoints, seed 0. SAMGPT: 500 updates, seed 39.

Within-model effects: GraphSAGE uses frozen linear probes; PRODIGY and SAMGPT use episodic few-shot evaluation. Budgets/readouts differ.